



Multi-Token Prediction (V3): denser signal per position, and it enables speculative decoding -- schematic

Standard: predict next token only



one target per position

MTP: sequential modules keep the full causal chain



each module reuses the previous depth's representation (not parallel independent heads)