

Lecture 17: Recurrent Neural Networks

Vanilla RNN · LSTM / GRU · Attention

Outline

Why Recurrent Networks?

BPTT and the Vanishing Gradient Problem

Gated Architectures: LSTM, GRU

Attention as Soft Lookup

Sequence Modeling: language, time series, speech